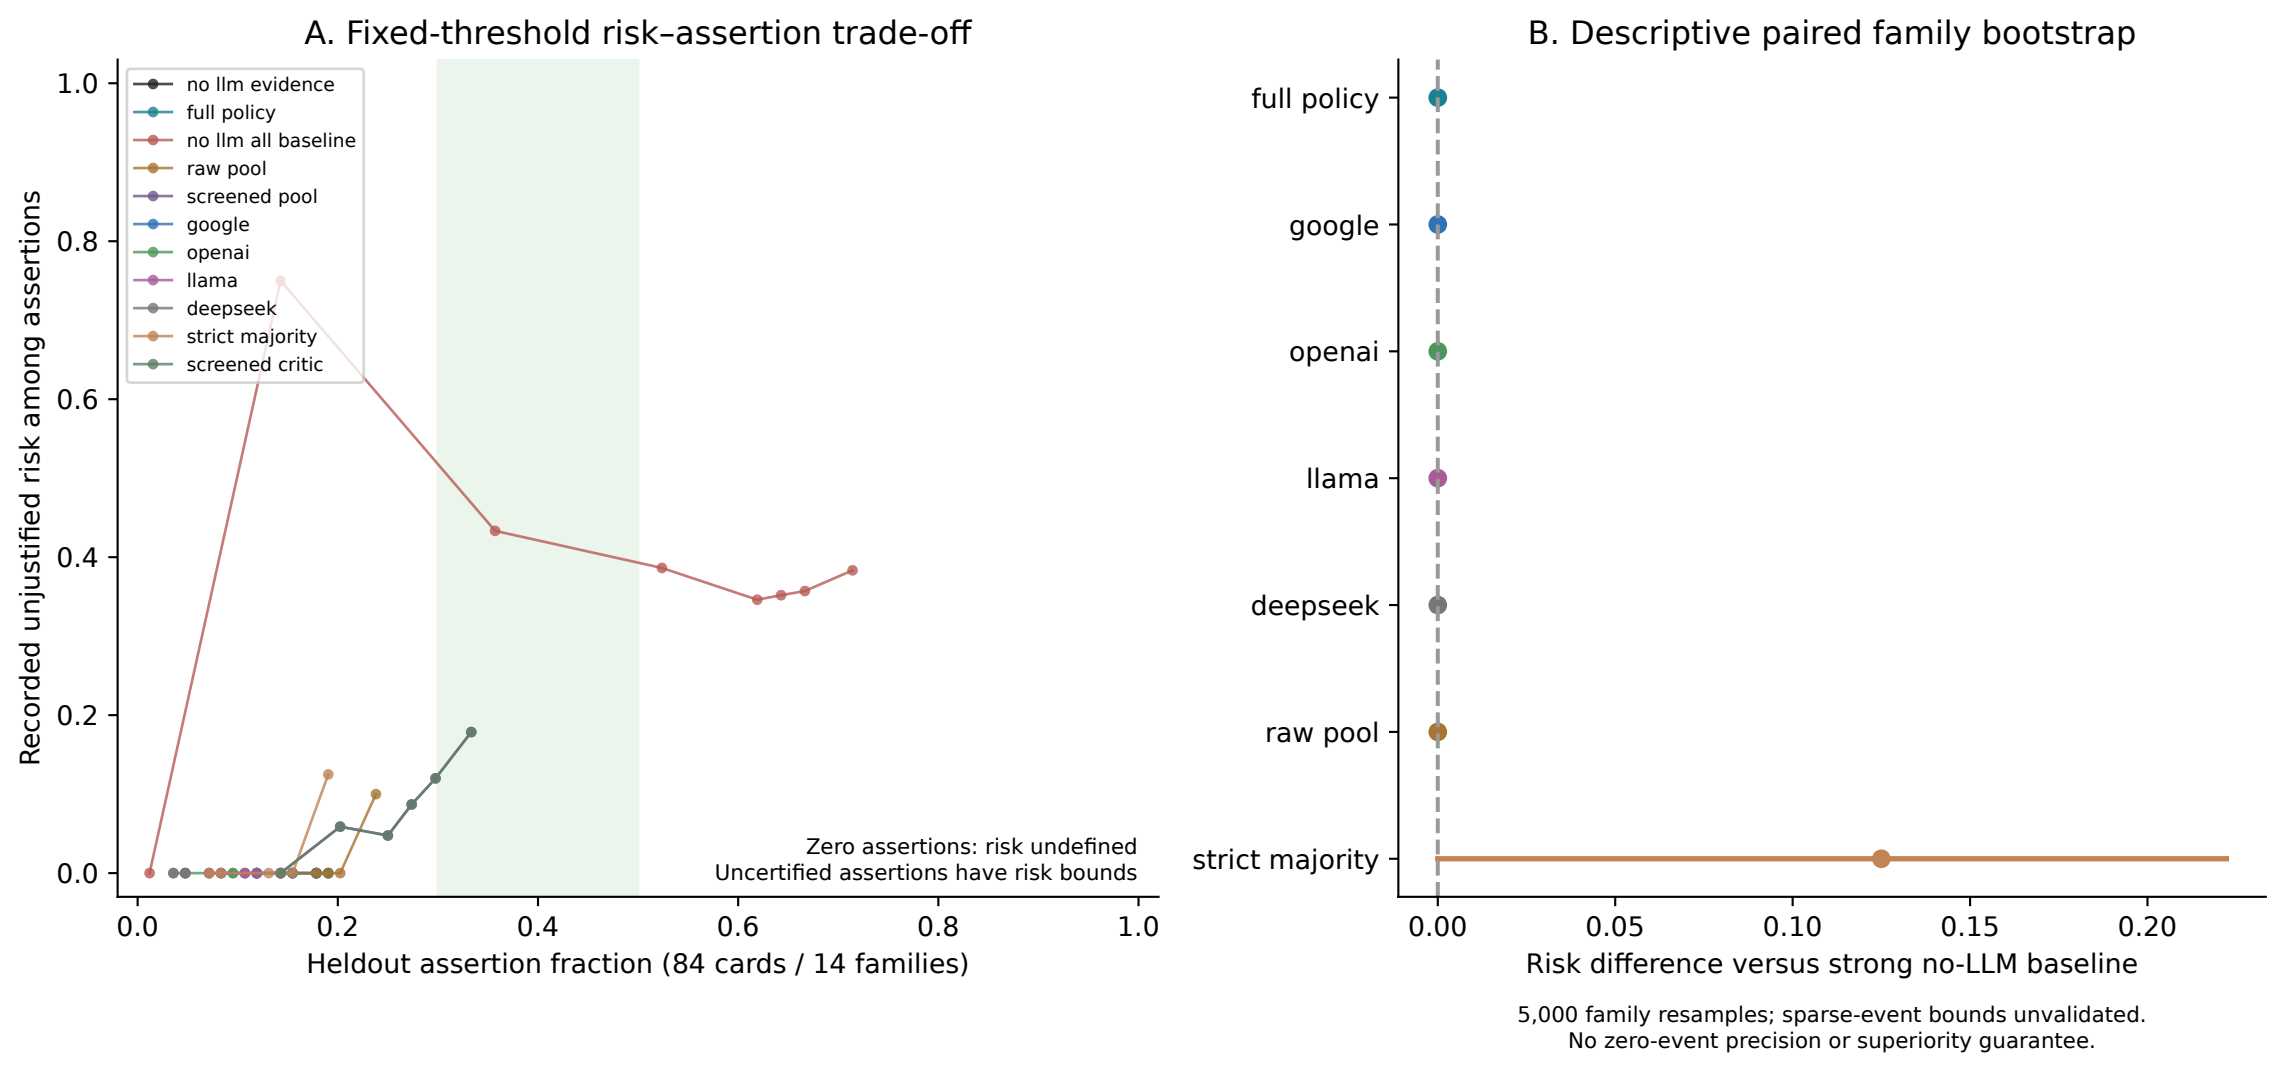
A. Fixed-threshold risk–assertion trade-off
Recorded unjustified risk among assertions
Heldout assertion fraction (84 cards / 14 families)
no llm evidence
full policy
no llm all baseline
raw pool
screened pool
google
openai
llama
deepseek
strict majority
screened critic
Zero assertions: risk undefined
Uncertified assertions have risk bounds
B. Descriptive paired family bootstrap
full policy
google
openai
llama
deepseek
raw pool
strict majority
Risk difference versus strong no-LLM baseline
5,000 family resamples; sparse-event bounds unvalidated.
No zero-event precision or superiority guarantee.